

Particle Filters

Stefan Contiu

- UTCN -

April 2015

Particle Filters

- Basic Algorithm
- Importance Sampling
- Mathematical Derivation of the PF
- Low Variance Resampling
- Robot Localization Example
- Conclusion

Particle Filters

- Alternative non-parametric implementation of Bayes filter.
- Key Idea : represent the posterior $bel(x_t)$ by a set of random state samples drawn from the posterior.

Call the samples : **Particles**, denote by:


$$X_t = x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]}$$

- each particle is a hypothesis as to what the true state may be at time t.

The Particles

- Ideally, the likelihood for a state hypothesis x_t to be included in the particle set X_t shall be proportional to its Bayes filter posterior $bel(x_t)$:

$$x_t^{[m]} \sim p(x_t \mid z_{1:t}, u_{1:t})$$



$$p(x_t \in X_t) \approx p(x_t \mid z_{1:t}) \quad (\text{equality for } M \uparrow \infty)$$

Basic Algorithm

```
1 : Algorithm Particle_fitter( $X_{t-1}, u_t, z_t$ ):  
2 :    $\overline{X}_t = X_t = \emptyset$   
3 :   for  $m = 1$  to  $M$  do  
4 :     sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$   
5 :      $w_t^{[m]} = p(z_t \mid x_t^{[m]})$   
6 :      $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7 :   endfor  
8 :   for  $m = 1$  to  $M$  do  
9 :     draw  $i$  with probability  $\propto w_t^{[i]}$   
10:    add  $x_t^{[i]}$  to  $X_t$   
11:  endfor  
12:  return  $X_t$ 
```

Basic Algorithm

1 : **Algorithm Particle_fitter**(X_{t-1}, u_t, z_t):

2 : $\overline{X}_t = X_t = \emptyset$

3 : for $m = 1$ to M do

4 : sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5 : $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6 : $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7 : endfor

8 : for $m = 1$ to M do

9 : draw i with probability $\propto w_t^{[i]}$

10: add $x_t^{[i]}$ to X_t

11: endfor

12: return X_t

Sequential importance sampling

Construct a temporary particle set \overline{X}_t reminiscent to the belief $\overline{bel}(x_t)$

Basic Algorithm

1 : **Algorithm Particle_fitter**(X_{t-1}, u_t, z_t):

2 : $\overline{X}_t = X_t = \emptyset$

3 : for $m = 1$ to M do

4 : sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5 : $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6 : $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7 : endfor

8 : for $m = 1$ to M do

9 : draw i with probability $\propto w_t^{[i]}$

10: add $x_t^{[i]}$ to X_t

11: endfor

12: return X_t

- Generates a hypothetical state x_t for time t based on particle x_{t-1} and u_t
- We need to sample from $p(x_t \mid u_t, x_{t-1})$. The ability of sample from the state transition probability is not given for arbitrary distributions. However many major distributions posses efficient algorithms for generating samples.
- Set of sampled particles is the representation of $\overline{bel}(x_t)$

Basic Algorithm

1 : **Algorithm Particle_fitter**(X_{t-1}, u_t, z_t):

2 : $\overline{X}_t = X_t = \emptyset$

3 : for $m = 1$ to M do

4 : sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5 : $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6 : $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7 : endfor

8 : for $m = 1$ to M do

9 : draw i with probability $\propto w_t^{[i]}$

10: add $x_t^{[i]}$ to X_t

11: endfor

12: return X_t

- Calculates for each particle $x_t^{[m]}$ the importance factor.
- Incorporates measurement z_t into the particle set.
- Set of weighted particles represents (in aprox) $bel(x_t)$

Basic Algorithm

1 : **Algorithm Particle_fitter**(X_{t-1}, u_t, z_t):

2 : $\overline{X}_t = X_t = \emptyset$

3 : for $m = 1$ to M do

4 : sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5 : $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6 : $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7 : endfor

8 : for $m = 1$ to M do

9 : draw i with probability $\propto w_t^{[i]}$

10: add $x_t^{[i]}$ to X_t

11: endfor

12: return X_t

Add the particle and its weight to the temporary particle set.

Basic Algorithm

1 : **Algorithm Particle_fitter**(X_{t-1}, u_t, z_t):

2 : $\overline{X}_t = X_t = \emptyset$

3 : for $m = 1$ to M do

4 : sample $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$

5 : $w_t^{[m]} = p(z_t \mid x_t^{[m]})$

6 : $\overline{X}_t = \overline{X}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$

7 : endfor

8 : for $m = 1$ to M do

9 : draw i with probability $\propto w_t^{[i]}$

10 : add $x_t^{[i]}$ to X_t


11 : endfor

12 : return X_t


Resampling:

- Eliminate particles with small importance weights
- Concentrate on particles with large weights

PF Algorithm Example

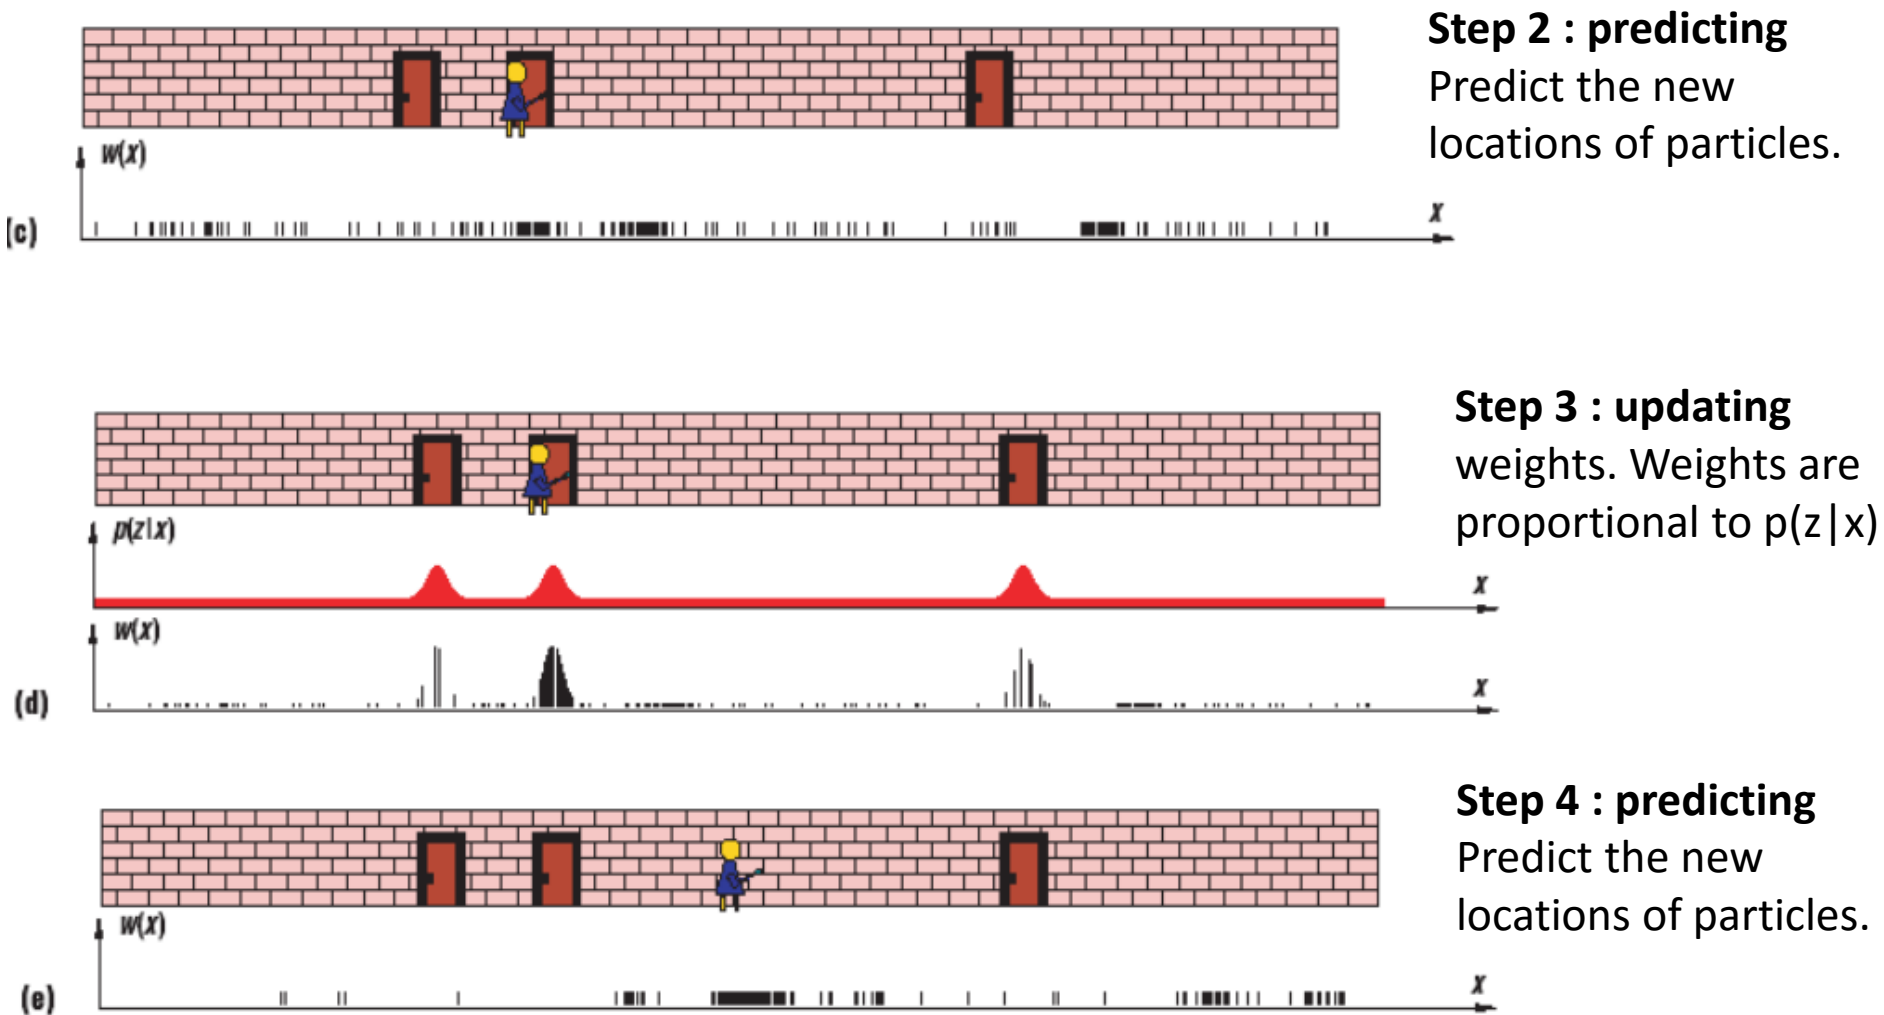


Step 0 : initialization
Each particle has the same weight



Step 1 : updating weights. Weights are proportional to $p(z|x)$

PF Algorithm Example



Importance Sampling

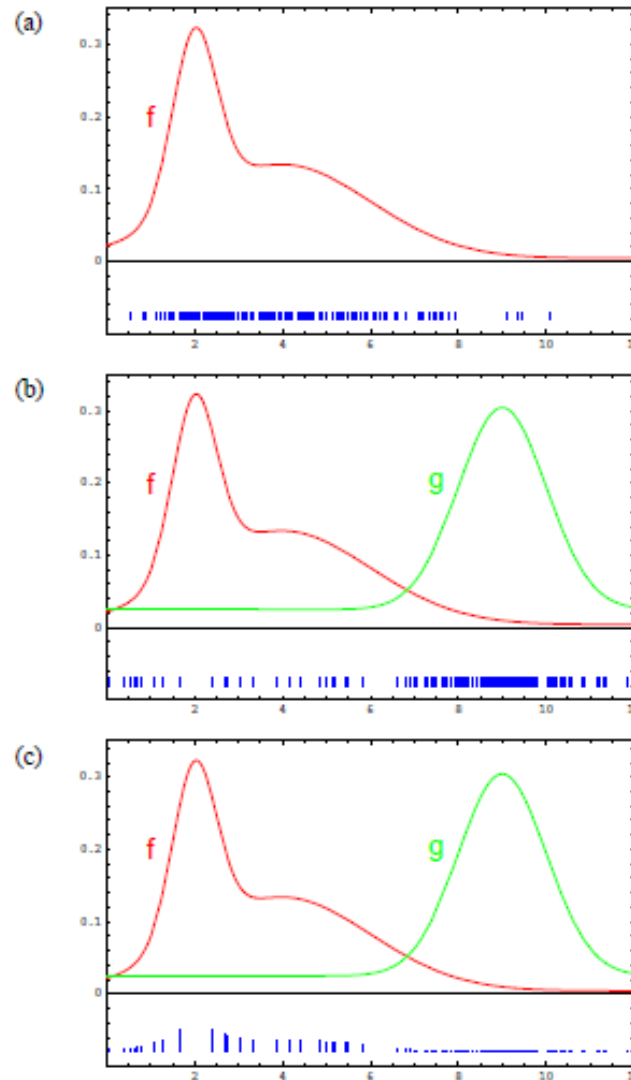
- Unfortunately it is often not possible to sample directly from the posterior distribution (f).
- We can use a different distribution g to generate samples from f .
- By introducing the importance weight w , we can account for the “differences between g and f ”.
- $w = f/g$
- f is often called target
- g is often called proposal

Importance Sampling

Samples cannot be drawn conveniently from the target distribution f .

Instead, the importance sampler draws samples from the proposal distribution g , which has a simpler form.

A sample of f is obtained by attaching the weight f/g to each sample x .



In particle filters:
 $f = \text{bel}(x_t)$
 $g = \overline{\text{bel}}(x_t)$

Mathematical Derivation of PF (1)

- Think of particles as samples of state sequences:

$$x_{0:t}^{[m]} = x_0^{[m]}, x_1^{[m]}, \dots, x_t^{[m]}$$

- Modify the algorithm:
 - Append to the particle $x_t^{[m]}$ the sequence of state samples from which it was generated $x_{0:t-1}^{[m]}$
 - This algorithm computes the posterior over all state sequences:

$$bel(x_{0:t}) = p(x_{0:t} \mid u_{1:t}, z_{1:t})$$

instead of the belief

$$bel(x_t) = p(x_t \mid u_{1:t}, z_{1:t})$$

Mathematical Derivation of PF (2)

- The posterior $bel(x_{0:t})$ is obtained in same way as the derivation of $bel(x_t)$ of Bayes Filters.
- First, we apply Bayes rule to the target posterior:

$$\begin{aligned}
 p(x_{0:t} \mid z_{1:t}, u_{1:t}) &\stackrel{\text{Bayes}}{=} \frac{p(z_t \mid x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} \mid z_{1:t-1}, u_{1:t})}{p(z_t \mid z_{1:t-1}, u_{1:t})} \\
 &= \eta p(z_t \mid x_{0:t}, z_{1:t-1}, u_{1:t}) p(x_{0:t} \mid z_{1:t-1}, u_{1:t})
 \end{aligned}$$

- Due to Conditional Independence:

$$\begin{aligned}
 &\stackrel{\text{Markov}}{=} \eta p(z_t \mid x_t) p(x_{0:t} \mid z_{1:t-1}, u_{1:t}) \\
 &= \eta p(z_t \mid x_t) p(x_t \mid x_{0:t-1}, z_{1:t-1}, u_{1:t}) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t}) \\
 &\stackrel{\text{Markov}}{=} \eta \underbrace{p(z_t \mid x_t)}_{\text{prediction}} p(x_t \mid x_{t-1}, u_t) \underbrace{p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})}_{\text{correction}}
 \end{aligned}$$

Mathematical Derivation of PF (3)

- Carry out derivation by induction.
- Initial Condition verification.
- For the mth particle:

$$\begin{aligned} p(x_t \mid x_{t-1}, u_t) \text{bel}(x_{0:t-1}) \\ = p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{0:t-1}, u_{0:t-1}) \end{aligned}$$

With

$$\begin{aligned} w_t^{[m]} &= \frac{\text{target distribution}}{\text{proposal distribution}} \\ &= \frac{\eta p(z_t \mid x_t) p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{1:t-1}, u_{1:t-1})}{p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{0:t-1}, u_{0:t-1})} \\ &= \eta p(z_t \mid x_t) \end{aligned}$$

Mathematical Derivation of PF (4)

- By resampling particles with probability proportional to w_t :


$$\eta w_t^{[m]} p(x_t \mid x_{t-1}, u_t) p(x_{0:t-1} \mid z_{0:t-1}, u_{0:t-1}) = \text{bel}(x_{0:t})$$

- The resulting particles are distributed according to the product of the proposal and the importance weight.
- This derivation is correct only for $M \rightarrow \infty$. However, even for finite M it explains the intuition behind PF.

Resampling


Given : Set X of wieghted samples.

Wanted : Random sample, where the probability of drawing x_i is given by w_i .



Roulette wheel

Generating n random values
 $O(n \lg n)$



Stochastic universal sampling

Generating one random value
 $O(n)$
Easy to implement, low variance


Low Variance Resampling

```
1:  Algorithm Low_variance_sampler( $\mathcal{X}_t, \mathcal{W}_t$ ):
2:     $\bar{\mathcal{X}}_t = \emptyset$ 
3:     $r = \text{rand}(0; M^{-1})$  ← Draw only once one random
4:     $c = w_t^{[1]}$                                      number  $r$  in  $(0, M^{-1})$ 
5:     $i = 1$ 
6:    for  $m = 1$  to  $M$  do
7:       $u = r + (m - 1) \cdot M^{-1}$  ← Add the fixed amount of  $M^{-1}$ 
8:      while  $u > c$                                      to  $r$ 
9:         $i = i + 1$ 
10:        $c = c + w_t^{[i]}$  ← Choose the corresponding
11:     endwhile                                           particle
12:     add  $x_t^{[i]}$  to  $\bar{\mathcal{X}}_t$ 
13:   endfor
14:   return  $\bar{\mathcal{X}}_t$ 
```


Advantages of Particle Filters

- can deal with non-linearities
- can deal with non-Gaussian noise
- mostly parallelizable
- easy to implement
- PFs focus adaptively on probable regions of state-space

Robot Localization using PF



Robot Localization using PF



References

- S. Thrun, W. Burgard, D. Fox: Probabilistic Robotics
- Stanford CS223B Computer Vision, Winter 2005, Lecture 12 Filters/Motion Tracking
- J. Xiao: Particle Filters, City College of New York.
- L. J. Latecki: Tutorial on Particle Filters, Temple University.